METRONOM

Introduction to Apache Calcite: working with data from any data source using SQL interface

Data sources

Introduction to Apache Calcite

Simplify writing of DBMS

Create a DBMS by composing

Advance the state of the art by pooling resources

Introduction to Apache Calcite

Connect datasource to Calcite

```
"schemas": [ {
    "name": "CSV",
    "type": "custom",
    "factory": "org.apache.calcite.adapter.csv.CsvSchemaFactory",
    "operand": {
      "directory": "csv_files"
}]
```

Connect datasource to Calcite

```
"name": "POSTGRESQL",
"type": "custom",
"factory": "org.apache.calcite.adapter.jdbc.JdbcSchema$Factory",
"operand": {
    "jdbcDriver": "org.postgresql.Driver",
    "jdbcUrl": "jdbc:postgresql://localhost:5432/postgres?...",
    "jdbcUser": "postgres",
    "jdbcPassword": "postgres"
}
```

Connect datasource to Calcite

```
"type": "custom",
"name": "ELASTIC",
"factory": "org.apache.calcite...ElasticsearchSchemaFactory",
"operand": {
   "coordinates": "{'127.0.0.1': 9200}",
   "index": ".monitoring-es-6-2019.08.20"
}
```

DEMO time

Use power of datasources

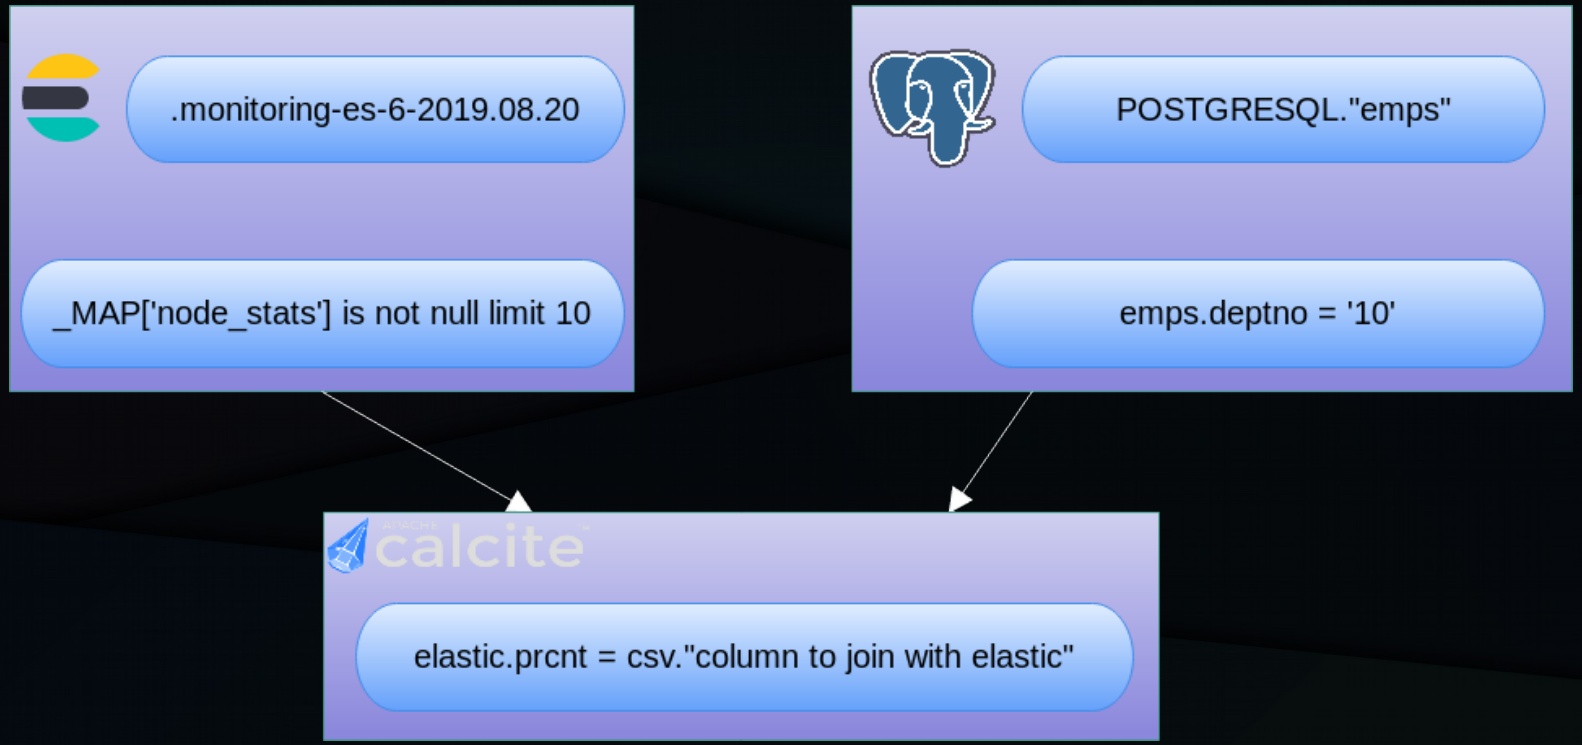

Powered by

Used by
APACHE STORM
beam
NGDATA
HIVE
APACHE DRILL
HerdDB
Apache Apex
druid
APACHE PHOENIX
Apache Solr
Alibaba Group
Apache Flink
samza
CASCADING
Qubole
dremio
omnisci
SQL Gremlin
UBER
Connects to
JDBC
APACHE Spark
splunk>
CSV
cassandra
{JSON}
elasticsearch
mongoDB
APACHE GEODE

Learn more

- https://github.com/snuyanzin/calcite-examples
- https://github.com/apache/calcite
- http://calcite.apache.org/
- http://calcite.apache.org/community/#mailing-lists